



Inspire...Educate...Transform.


Ensemble Learning

Dr. Manoj Chinnakotla

Senior Applied Scientist, Microsoft
Adjunct Professor, IIIT Hyderabad

18th March, 2017

What is an Ensemble?



**Combine multiple
models into one!**



Applications: classification, clustering,
collaborative filtering, anomaly detection.....

Stories of Success



- **Million-Dollar Prize**

- Improve the baseline movie recommendation approach of Netflix by 10% in accuracy
- The top submissions all combine several teams and algorithms as an ensemble



- **Data Mining Competitions**

- Classification problems
- Winning teams employ an ensemble of classifiers

Netflix Prize

- **Supervised Learning Task**

- Training data is a set of users and ratings (1,2,3,4,5 stars) those users have given to movies.
- Construct a classifier that given a user and an unrated movie, correctly classifies that movie as either 1, 2, 3, 4, or 5 stars
- \$1 million prize for a 10% improvement over Netflix's current movie recommender

- **Competition**

- At first, single-model methods are developed, and performances are improved
- However, improvements slowed down
- Later, individuals and teams merged their results, and significant improvements are observed

Leaderboard

Rank	Team Name	Best Test Score	% Improvement	Best Submit Time
Grand Prize - RMSE = 0.8567 - Winning Team: BellKor's Pragmatic Chaos				
1	BellKor's Pragmatic Chaos	0.8567	10.06	2009-07-26 18:18:28
2	The Ensemble	0.8567	10.06	2009-07-26 18:38:22
3	Grand Prize Team	0.8582	9.90	2009-07-10 21:24:40
4	Opera Solutions and Vandelay United	0.8588	9.84	2009-07-10 01:12:31
5	Vandelay Industries !	0.8591	9.81	2009-07-10 00:32:20
6	PragmaticTheory	0.8594	9.77	2009-06-24 12:06:56
7	BellKor in BigChaos	0.8601	9.70	2009-05-13 08:14:09
8	Dace	0.8612	9.59	2009-07-24 17:18:43
9	Feeds2	0.8622	9.48	2009-07-12 13:11:51
10	BigChaos	0.8623	9.47	2009-04-07 12:33:59

“Our final solution (RMSE=0.8567) consists of blending 107 individual results. “

12	BellKor	0.8624	9.46	2009-07-26 17:19:11
Progress Prize 2008 - RMSE = 0.8627 - Winning Team: BellKor in BigChaos				
13	xianqiang	0.8642	9.27	2009-07-15 14:53:22
14	Gravity	0.8643	9.26	2009-04-22 18:31:32
15	Ces	0.8651	9.18	2009-06-21 19:24:53

“Predictive accuracy is substantially improved when blending multiple predictors. Our experience is that most efforts should be concentrated in deriving substantially different approaches, rather than refining a single technique. “

Progress Prize 2007 - RMSE = 0.8723 - Winning Team: Korben

Cinematix score - RMSE = 0.9525

Motivation for Ensembles

- Ensemble model improves accuracy and robustness over single model methods
- Applications:
 - distributed computing
 - privacy-preserving applications
 - large-scale data with reusable models
 - multiple sources of data
- Efficiency: a complex problem can be decomposed into multiple sub-problems that are easier to understand and solve (divide-and-conquer approach)

Why Ensemble Works? (1)

- **Intuition**


- Combining diverse, independent opinions in human decision-making as a protective mechanism (e.g. stock portfolio)

- **Uncorrelated Error Reduction**

- Suppose we have 5 completely independent classifiers for majority voting
- If accuracy is 70% for each
 - $10 (.7^3)(.3^2)+5(.7^4)(.3)+(.7^5)$
 - **83.7% majority vote accuracy**
- 101 such classifiers
 - **99.9% majority vote accuracy**

from T. Holloway, Introduction to Ensemble Learning, 2007.


Why Ensemble Works? (2)




Ensemble gives the global picture!

Why Ensemble Works? (3)

- Overcome limitations of single hypothesis
 - The target function may not be implementable with individual classifiers, but may be approximated by model averaging



Decision Tree




Model Averaging

Generalization Error & Bias, Variance Tradeoff

Bias and Variance

- Bias
 - Measures the accuracy or quality of the algorithm
 - High bias means a poor match
- Variance
 - Measures the precision or specificity of the match
 - A high variance means a weak match
- We would like to minimize each of these
- Unfortunately, we can't do this independently, since there is a trade-off

Bias and Variance (Contd..)



Bias-Variance Analysis in Regression


True function is $y = f(x) + e$

- Where e is normally distributed with zero mean and standard deviation s
- Given a set of training examples, $\{(x_i, y_i)\}$, we fit an hypothesis $h(x) = w \cdot x + b$ to the data to minimize the squared error

$$\sum_i (y_i - h(x_i))^2$$


Bias-Variance Analysis in Regression (Contd..)

$$y = x + 2 \sin(1.5x) + N(0,0.2)$$



Bias-Variance Analysis in Regression (Contd..)

50 fits (20 examples each)



Bias-Variance Analysis in Regression

(Contd..)

- Given a new point x^* , with observed value $y^* = f(x^*) + \varepsilon$, let's estimate the expected prediction error given by:

$$E[(y^* - h(x^*))^2]$$

Bias-Variance Analysis in Regression (Contd..)

- Imagine that our training sample S is drawn from a population with distribution $P(S)$ then we need to compute:

$$E_P[(y^* - h(x^*))^2]$$

Bias-Variance Analysis in Regression (Contd..)

- Let Z be a random variable with probability distribution $P(Z)$
- Let $\bar{Z} = E_p[Z]$ be the average value of Z .
- Lemma: $E[(Z - \bar{Z})^2] = E[Z^2] - \bar{Z}^2$

$$\begin{aligned} E[(Z - \bar{Z})^2] &= E[Z^2 - 2 Z \bar{Z} + \bar{Z}^2] \\ &= E[Z^2] - 2 E[Z] \bar{Z} + \bar{Z}^2 \\ &= E[Z^2] - 2 \bar{Z}^2 + \bar{Z}^2 \\ &= E[Z^2] - \bar{Z}^2 \end{aligned}$$

- Corollary : $E[Z^2] = E[(Z - \bar{Z})^2] + \bar{Z}^2$

Bias-Variance Analysis in Regression (Contd..)





$$\begin{aligned}
 E[(h(x^*) - y^*)^2] &= E[h(x^*)^2 - 2h(x^*)y^* + y^{*2}] \\
 &= E[h(x^*)^2] - 2E[h(x^*)]E[y^*] + E[y^{*2}] \\
 &= E[(h(x^*) - \overline{h(x^*)})^2] + \overline{h(x^*)}^2 \\
 &\quad - 2\overline{h(x^*)}f(x^*) \\
 &\quad + E[(y^* - f(x^*))^2] + f(x^*)^2 \\
 &= E[(h(x^*) - \overline{h(x^*)})^2] + \text{VARIANCE} \\
 &\quad \left(\overline{h(x^*)}^2 - f(x^*) \right)^2 \text{BIAS} \\
 &\quad + E[(y^* - f(x^*))^2] \text{NOISE} \\
 &= \text{Var}(h(x^*)) + \text{Bias}(h(x^*))^2 + E[\varepsilon^2] \\
 &= \text{Var}(h(x^*)) + \text{Bias}(h(x^*))^2 + \sigma^2
 \end{aligned}$$

Expected prediction error = Variance + Bias² + Noise²



Bias-Variance Analysis in Regression (Contd..)

- Variance: $E_P \left[\left(h(x^*) - \overline{h(x^*)} \right)^2 \right]$
 - Describes how much $h(x^*)$ varies from one training set S to another
- Bias: $\left[\overline{h(x^*)} - f(x^*) \right]$
 - Describes the average error of $h(x^*)$
- Noise: $E_P \left[(y^* - f(x^*))^2 \right]$
 - Describes how much y^* varies from the true function $f(x^*)$

Bias-Variance Analysis in Regression (Contd..)



Bias vs. Variance Tradeoff



Bagging

- Create ensembles by “*bootstrap aggregation*”, i.e., repeatedly randomly resampling the training data (Brieman, 1996).
 - Bootstrap: draw N items from X with replacement
- Bagging
 - Train M learners on M bootstrap samples
 - Combine outputs by voting (e.g., *majority vote*)
- Decreases error by *decreasing the variance* in the results due to *unstable learners*, algorithms (like decision trees and neural networks) whose output can change dramatically when the training data is slightly changed.

Bagging – Aggregate Bootstrapping

- Given a standard training set D of size n
- For $i = 1 \dots M$
 - Draw a sample of size $n^* < n$ from D uniformly and with replacement
 - Learn classifier C_i
- Final classifier is a vote of $C_1 \dots C_M$
- Increases classifier stability/reduces variance

Bagging – An Example

Original Data:

x	0.1	0.2	0.3	0.4	0.5	0.6	0.7	0.8	0.9	1
y	1	1	1	-1	-1	-1	-1	1	1	1

Bootstrap samples and classifiers:

x	0.1	0.2	0.2	0.3	0.4	0.4	0.5	0.6	0.9	0.9
y	1	1	1	1	-1	-1	-1	-1	1	1

x	0.1	0.2	0.3	0.4	0.5	0.5	0.9	1	1	1
y	1	1	1	-1	-1	-1	1	1	1	1

x	0.1	0.2	0.3	0.4	0.4	0.5	0.7	0.7	0.8	0.9
y	1	1	1	-1	-1	-1	-1	-1	1	1

x	0.1	0.2	0.5	0.5	0.5	0.7	0.7	0.8	0.9	1
y	1	1	-1	-1	-1	-1	-1	1	1	1

Combine predictions by majority voting

Random Forests*

- **Algorithm**
 - Choose T —number of trees to grow
 - Choose $m < M$ (M is the number of total features) —number of features used to calculate the best split at each node (typically 20%)
 - For each tree
 - Choose a training set by choosing N times (N is the number of training examples) with replacement from the training set
 - For each node, randomly choose m features and calculate the best split
 - Fully grown and not pruned
 - Use majority voting among all the trees

*[Breiman01]

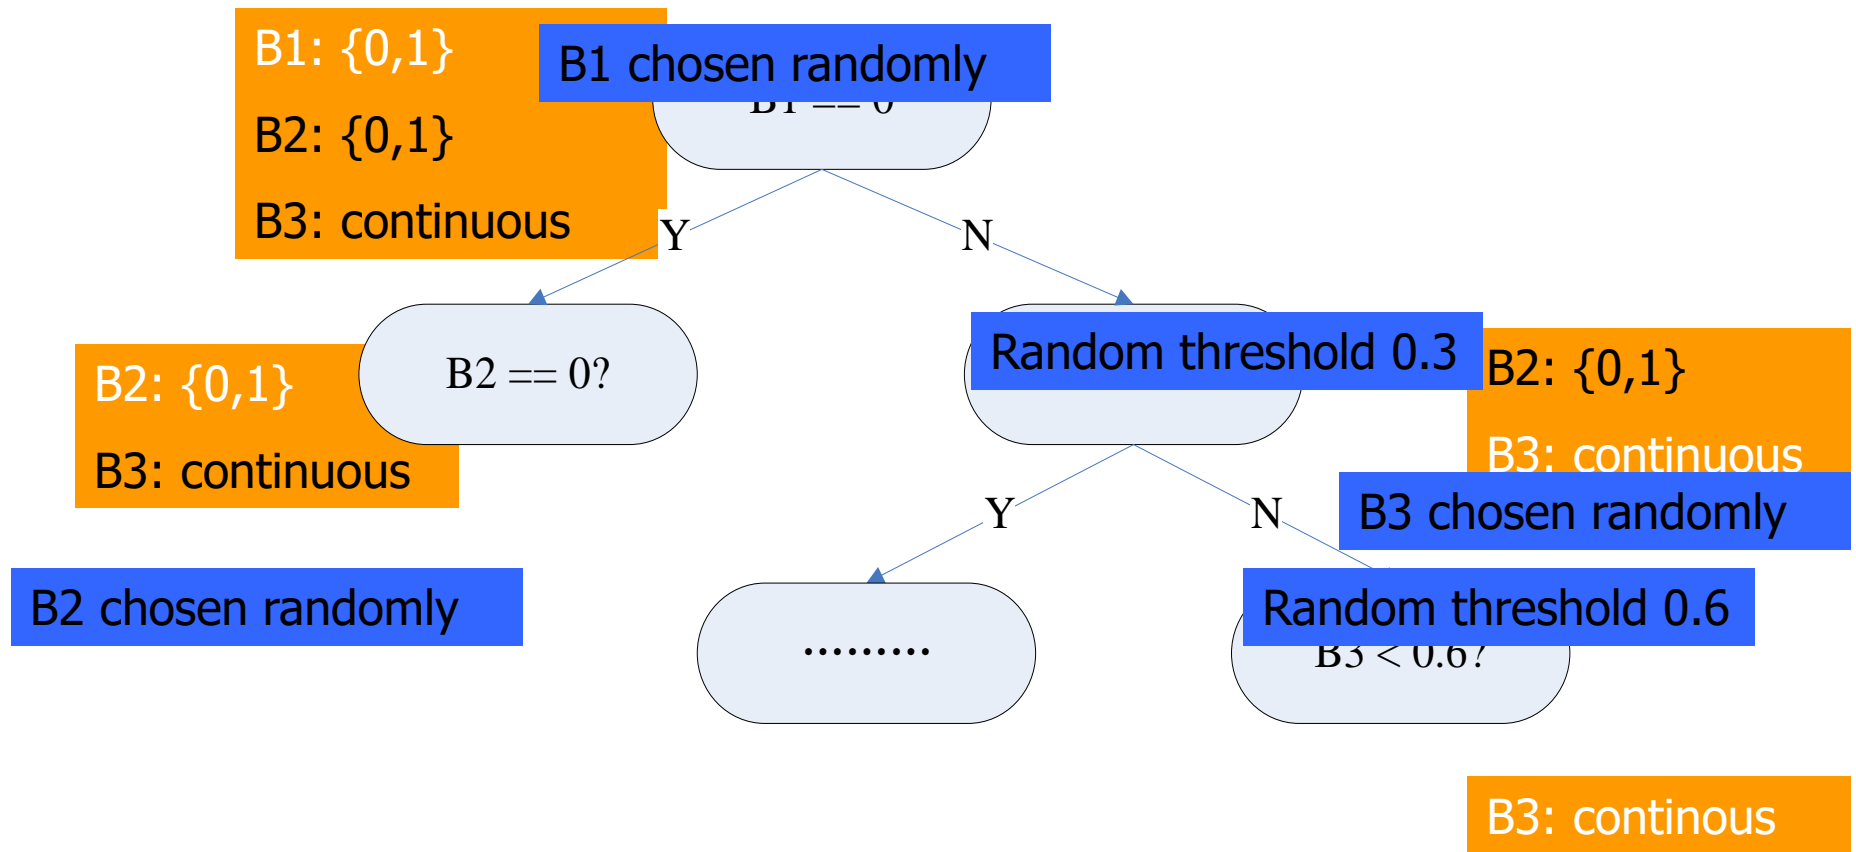
Random Forests

- Discussions
 - Bagging + random features
 - Improve accuracy
 - Incorporate more diversity and reduce variances
 - Improve efficiency
 - Searching among subsets of features is much faster than searching among the complete set

Random Decision Tree

- **Algorithm**
 - At each node, an un-used feature is chosen randomly
 - A discrete feature is un-used if it has never been chosen previously on a given decision path starting from the root to the current node.
 - A continuous feature can be chosen multiple times on the same decision path, but each time a different threshold value is chosen
 - We stop when one of the following happens:
 - A node becomes too small (≤ 3 examples).
 - Or the total height of the tree exceeds some limits, such as the total number of features.
 - Prediction
 - Simple averaging over multiple trees

Random Decision Tree




Random Decision Tree


- **Potential Advantages**
 - Training can be very efficient. Particularly true for very large datasets.
 - No cross-validation based estimation of parameters for some parametric methods.
 - Natural multi-class probability.
 - Imposes very little about the structures of the model.

Optimal Decision Boundary

Figure 3.5: Gaussian mixture training samples and optimal boundary.




training samples




optimal boundary



from Tony Liu's thesis (supervised by Kai Ming Ting)




(a) unpruned C4.5



(b) Bagging



(c) Random Forests



(d) Complete-random tree ensemble

RDT looks like the optimal boundary

Strong vs. Weak Learners

- **Strong Learner** → Objective of machine learning
 - Take labeled data for training
 - Produce a classifier which can be *arbitrarily accurate*
- **Weak Learner**
 - Take labeled data for training
 - Produce a classifier which is *more accurate than random guessing*

Boosting – Combine Weak Learners

- **Weak Learner:** only needs to generate a hypothesis with a training accuracy greater than 0.5, i.e., < 50% error over any distribution
- Learners
 - Strong learners are very difficult to construct
 - Constructing weaker Learners is relatively easy
- Questions: Can a set of **weak learners** create a single **strong learner** ?
- YES 😊

Boost weak classifiers to a strong learner

Boosting*

- **Principles**

- Boost a set of weak learners to a strong learner
- Make records currently misclassified more important

- **Example**

- Record 4 is hard to classify
- Its weight is increased, therefore it is more likely to be chosen again in subsequent rounds

Original Data	1	2	3	4	5	6	7	8	9	10
Boosting (Round 1)	7	3	2	8	7	9	4	10	6	3
Boosting (Round 2)	5	4	9	4	2	5	1	7	4	2
Boosting (Round 3)	4	4	8	10	4	5	4	6	3	4

*[FrSc97]

from P. Tan et al. Introduction to Data Mining.

Boosting

- AdaBoost
 - Initially, set uniform weights on all the records
 - At each round
 - Create a bootstrap sample based on the weights
 - Train a classifier on the sample and apply it on the original training set
 - Records that are wrongly classified will have their weights increased
 - Records that are classified correctly will have their weights decreased
 - If the error rate is higher than 50%, start over
 - Final prediction is weighted average of all the classifiers with weight representing the training accuracy

Boosting

- Determine the weight

- For classifier i , its error is
- The classifier's importance is represented as:
- The weight of each record is updated as:
- Final combination:


$$\varepsilon_i = \frac{\sum_{j=1}^N w_j \delta(C_i(x_j) \neq y_j)}{\sum_{j=1}^N w_j}$$

$$\alpha_i = \frac{1}{2} \ln \left(\frac{1 - \varepsilon_i}{\varepsilon_i} \right)$$


$$w_j^{(i+1)} = \frac{w_j^{(i)} \exp(-\alpha_i y_j C_i(x_j))}{Z^{(i)}}$$

$$C^*(x) = \arg \max_y \sum_{i=1}^K \alpha_i \delta(C_i(x) = y)$$

Boosting – An Example



AdaBoost Animation




Original Training Set: Equal weights to all training samples

Taken from “A Tutorial on Boosting” by Yoav Freund and Rob Schapire

AdaBoost Animation


ε = error rate of classifier
 α = weight of classifier



$\varepsilon_1 = 0.30$
 $\alpha_1 = 0.42$




D_2




AdaBoost Animation

ROUND 2




$$\epsilon_2 = 0.21$$
$$\alpha_2 = 0.65$$



AdaBoost Animation

ROUND 3



AdaBoost Animation

$$H_{\text{final}} = \text{sign} \left(0.42 \begin{array}{|c|} \hline \text{blue} \\ \hline \text{red} \end{array} + 0.65 \begin{array}{|c|} \hline \text{blue} \\ \hline \text{red} \end{array} + 0.92 \begin{array}{|c|} \hline \text{blue} \\ \hline \text{red} \end{array} \right)$$


The diagram shows three weak classifiers, each represented by a square divided into a blue region and a red region by a vertical line. The first classifier has a weight of 0.42, the second has a weight of 0.65, and the third has a weight of 0.92. These are summed and passed through a sign function to produce the final hypothesis H_{final} .

Mixture of Experts

- Voting where weights are input-dependent (gating)
- Different input regions covered by different learners (Jacobs et al., 1991)


$$y = \sum_{j=1}^L w_j d_j$$

- Gating decides which expert to use
- Need to learn the individual experts as well as the gating functions $w_i(x)$:
 $\sum w_j(x) = 1$, for all x



Stacking

- Combiner $f()$ is another learner (Wolpert, 1992)



Strengths of AdaBoost

- It has no parameters to tune (except for the number of rounds)
- It is fast, simple and easy to program (relatively)
- It comes with a set of theoretical guarantee (e.g., training error, test error)
- Instead of trying to design a learning algorithm that is accurate over the entire space, we can focus on finding base learning algorithms that only need to be better than random.
- It can identify outliers: i.e. examples that are either mislabeled or that are inherently ambiguous and hard to categorize.

Weakness of AdaBoost

- The actual performance of boosting depends on the data and the base learner.
- Boosting seems to be especially susceptible to noise.
- When the number of outliers is very large, the emphasis placed on the hard examples can hurt the performance.
➔ “Gentle AdaBoost”, “BrownBoost”

References

- *Survey of Boosting from an Optimization Perspective*. Manfred K. Warmuth and S.V.N. Vishwanathan. ICML'09, Montreal, Canada, June 2009.
- *Theory and Applications of Boosting*. Robert Schapire. NIPS'07, Vancouver, Canada, December 2007.
- *From Trees to Forests and Rule Sets--A Unified Overview of Ensemble Methods*. Giovanni Seni and John Elder. KDD'07, San Jose, CA, August 2007.

International School of Engineering

Plot 63/A, 1st Floor, Road # 13, Film Nagar, Jubilee Hills, Hyderabad - 500 033

For Individuals: +91-9502334561/63 or 040-65743991

For Corporates: +91-9618483483

Web: <http://www.insofe.edu.in>

Facebook: <https://www.facebook.com/insofe>

Twitter: <https://twitter.com/Insofeedu>

YouTube: <http://www.youtube.com/InsofeVideos>

SlideShare: <http://www.slideshare.net/INSOFE>

LinkedIn: <http://www.linkedin.com/company/international-school-of-engineering>

This presentation may contain references to findings of various reports available in the public domain. INSOF makes no representation as to their accuracy or that the organization subscribes to those findings.